

INSURANCE DATASET

Logistic Regression

1.PROBLEM STATEMENT: To predict and analyze the Female and Male Smoker in the region

```
In [3]: import pandas as pd
import numpy as np
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
```

DATA COLLECTION

```
In [4]: df=pd.read_csv(r"C:\Users\manis\OneDrive\Pictures\Documents\insurance.csv")
df
```

```
Out[4]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

DATA CLEANING

```
In [3]: df.info()
#TO FIND IF THE NULL VALUE ARE HAVE OR NOT
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

```
In [4]: df['region'].value_counts()
#TO KNOW THE REGION COUNT VALUE
```

```
Out[4]: region
southeast    364
southwest    325
northwest    325
northeast    324
Name: count, dtype: int64
```

```
In [5]: convert={'sex':{'female':1,"male":0}}
df=df.replace(convert)
df
#REPLACED THE FEMALE AS 1 AND MALE AS 0 BECAUSE IT WAS IN STRING
```

```
Out[5]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	yes	southwest	16884.92400
1	18	0	33.770	1	no	southeast	1725.55230
2	28	0	33.000	3	no	southeast	4449.46200
3	33	0	22.705	0	no	northwest	21984.47061
4	32	0	28.880	0	no	northwest	3866.85520
...
1333	50	0	30.970	3	no	northwest	10600.54830
1334	18	1	31.920	0	no	northeast	2205.98080
1335	18	1	36.850	0	no	southeast	1629.83350
1336	21	1	25.800	0	no	southwest	2007.94500
1337	61	1	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [6]: convert={'region':{'southeast':1,"southwest":2,"northwest":3,"northeast":4}}
df=df.replace(convert)
df
#REPLACING THE STRING TO NUMERIC VALUES
```

Out[6]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	yes	2	16884.92400
1	18	0	33.770	1	no	1	1725.55230
2	28	0	33.000	3	no	1	4449.46200
3	33	0	22.705	0	no	3	21984.47061
4	32	0	28.880	0	no	3	3866.85520
...
1333	50	0	30.970	3	no	3	10600.54830
1334	18	1	31.920	0	no	4	2205.98080
1335	18	1	36.850	0	no	1	1629.83350
1336	21	1	25.800	0	no	2	2007.94500
1337	61	1	29.070	0	yes	3	29141.36030

1338 rows × 7 columns

```
In [7]: x=['sex','bmi','children','region','charges']
        y=["yes","no"]
```

```
In [8]: all_inputs=df[x]
        all_classes=df['smoker']
        x_train,x_test,y_train,y_test=train_test_split(all_inputs,all_classes,train_size=0.8)
```

```
In [9]: dc=DecisionTreeClassifier()
        dc.fit(x_train,y_train)
        #FITTING THE X AND TRAIN IN THE DECISION TREE CLASSIFIER
```

```
Out[9]: ▾ DecisionTreeClassifier
        DecisionTreeClassifier()
```

```
In [10]: dc.score(x_test,y_test)
         #TO FIND THE ACCURACY SCORE
```

Out[10]: 0.9387144992526159

RANDOM FOREST

using insurance dataset

```
In [11]: from sklearn.ensemble import RandomForestClassifier
         #LIBRARY FOR RANDOMFOREST
```

```
In [12]: rf=RandomForestClassifier()
         rf.fit(x_train,y_train)
         #TO FIT THE X AND Y TRAIN OF RANDOMFOREST
```

Out[12]: ▾ RandomForestClassifier
RandomForestClassifier()

```
In [13]: rf=RandomForestClassifier()
params={'max_depth':[2,3,4,5,6], 'min_samples_leaf':[5,10,15,20,50,100], 'n_estimators':100}
#PARAMETERS ARE USED TO SPLIT NODES
```

```
In [14]: from sklearn.model_selection import GridSearchCV
#GRIDSEARCH IS TO FIND THE PARAMETER VALUES FROM THE SET OF PARAMETERS THAT WERE
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring='accuracy')
grid_search.fit(x_train,y_train)
#TO FIT THE X AND Y TRAIN IN GRID SEARCH
```

Out[14]: ▸ **GridSearchCV**
▸ **estimator: RandomForestClassifier**
 ▸ RandomForestClassifier

```
In [15]: grid_search.best_score_
```

Out[15]: 0.9447269639824828

```
In [16]: rf_best=grid_search.best_estimator_
print(rf_best)
#THE ESTIMATOR IS STORED IN RF_BEST
#BEST ESTIMATOR IS USED TO CALL THE PREDICT AND SCORE

RandomForestClassifier(max_depth=6, min_samples_leaf=5)
```

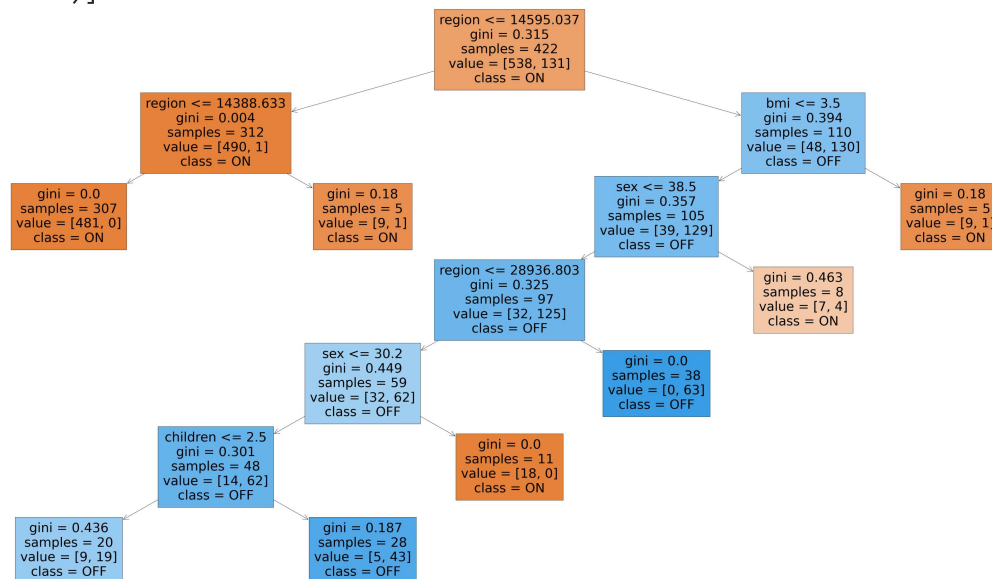
```
In [17]: x=df.drop('smoker',axis=1)
y=df['smoker']
#HERE THE SMOKER COLUMN WAS DROP IN X AND STORED IN Y
```

```
In [18]: from sklearn.tree import plot_tree
from sklearn.tree import DecisionTreeClassifier
import matplotlib.pyplot as plt
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[5],feature_names=x.columns,class_names=['ON','OFF'])
#TO PLOT THE DECISIONTREE WE NEED TO IMPORT THE LIBRARY PLOT_TREE AND DECISIONTREE
```

```

Out[18]: [Text(0.5, 0.9285714285714286, 'region <= 14595.037\ngini = 0.315\nsamples = 422\nvalue = [538, 131]\nnclass = ON'),
  Text(0.25, 0.7857142857142857, 'region <= 14388.633\ngini = 0.004\nsamples = 312\nvalue = [490, 1]\nnclass = ON'),
  Text(0.125, 0.6428571428571429, 'gini = 0.0\nsamples = 307\nvalue = [481, 0]\nnclass = ON'),
  Text(0.375, 0.6428571428571429, 'gini = 0.18\nsamples = 5\nvalue = [9, 1]\nnclass = ON'),
  Text(0.75, 0.7857142857142857, 'bmi <= 3.5\ngini = 0.394\nsamples = 110\nvalue = [48, 130]\nnclass = OFF'),
  Text(0.625, 0.6428571428571429, 'sex <= 38.5\ngini = 0.357\nsamples = 105\nvalue = [39, 129]\nnclass = OFF'),
  Text(0.5, 0.5, 'region <= 28936.803\ngini = 0.325\nsamples = 97\nvalue = [32, 125]\nnclass = OFF'),
  Text(0.375, 0.35714285714285715, 'sex <= 30.2\ngini = 0.449\nsamples = 59\nvalue = [32, 62]\nnclass = OFF'),
  Text(0.25, 0.21428571428571427, 'children <= 2.5\ngini = 0.301\nsamples = 48\nvalue = [14, 62]\nnclass = OFF'),
  Text(0.125, 0.07142857142857142, 'gini = 0.436\nsamples = 20\nvalue = [9, 19]\nnclass = OFF'),
  Text(0.375, 0.07142857142857142, 'gini = 0.187\nsamples = 28\nvalue = [5, 43]\nnclass = OFF'),
  Text(0.5, 0.21428571428571427, 'gini = 0.0\nsamples = 11\nvalue = [18, 0]\nnclass = ON'),
  Text(0.625, 0.35714285714285715, 'gini = 0.0\nsamples = 38\nvalue = [0, 63]\nnclass = OFF'),
  Text(0.75, 0.5, 'gini = 0.463\nsamples = 8\nvalue = [7, 4]\nnclass = ON'),
  Text(0.875, 0.6428571428571429, 'gini = 0.18\nsamples = 5\nvalue = [9, 1]\nnclass = ON')]

```



```
In [19]: rf_best.feature_importances_
```

```
Out[19]: array([0.00629788, 0.08285052, 0.01589012, 0.0140862 , 0.88087528])
```

```
In [20]: df1=pd.DataFrame({'Varname':x_train.columns,'Imp':rf_best.feature_importances_})
```

```
In [21]: df1.sort_values(by='Imp',ascending=False)
```

Out[21]:

	Varname	Imp
4	charges	0.880875
1	bmi	0.082851
2	children	0.015890
3	region	0.014086
0	sex	0.006298


```
In [10]: v=pd.crosstab(df['smoker'],df['sex'])
v
```

Out[10]:

	sex	female	male
smoker			
no		547	517
yes		115	159

```
In [16]: v.plot(kind='bar',stacked=True,color=["lavender","grey"],grid=False)
```

Out[16]: <Axes: xlabel='smoker'>



CONCLUSION

IN DECISION TREE THE SCORE OF X AND Y IS 94% AND IN THE RANDOM FOREST THE SCORE IS 95% COMPARING THE BOTH RANDOM FOREST IS HIGHEST IN THE ACCURACY AND AS PER THE PROBLEM STATEMENT MALE SMOKERS ARE HIGHER THEN FEMALE SMOKER.